

An Overview of Neural Networks

Harley Clifton, Becky Catlett, Natasha Gesker, and Eliot Liucci

2023-10-10

Contents

Introduction	2
Deep Learning (Becky)	2
The Neural Network Model (Eliot)	2
The Input Layer	2
Nonlinear Activation Functions	2
Hidden Layers	2
Output Layer	3
Training Neural Networks	3
Applications of Principal Component Analysis with Neural Networks	5
Single-Layer Neural Networks (Harley)	5
Example	5
Multi-Layer Neural Networks (Natasha)	5
Example	5
References	6

Introduction

Deep Learning (Becky)

The Neural Network Model (Eliot)

A neural network is a type of deep learning algorithm that makes use of a web of nodes to predict or classify data. The complexity of a neural network can vary greatly based on what task is required. As the name suggests, they are similar in function to neurons in a brain. The model takes in information through the *input layer*, which then activates various nodes in the *hidden layers*, and then a result is produced.

The Input Layer

The input layer is where data can be input into the model. If we have p input variables, which we will denote $X = X_1, X_2, \dots, X_p$, then our network will have p input nodes. Each node in future layers will depend on the value that X_i holds.

Nonlinear Activation Functions

Before we get into the hidden layer, it is important to understand what is happening at each hidden layer node. Each hidden layer node is computed by taking a weighted linear combination of the input layer and then applying a *nonlinear activation function* so that the *activation*, which is the value the node will take based on input vector X , will be between 0 and 1.

We will discuss two of the most common activation functions. For simpler networks, the *sigmoid* function is effective. The sigmoid function is defined as

$$S(x) = \frac{1}{1 + e^{-x}}$$

As discussed previously, the purpose of the activation function is to bring the range of values for the input layer down to any value between 0 and 1.

Another activation function that is more common in networks that require more “training” is the rectified linear activation unit function, or ReLU for short. The ReLU function is defined as

$$R(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{otherwise} \end{cases}$$

The benefits of using ReLU over Sigmoid is that ReLU can be better used for *backpropagation*, which is the main technique used to train networks. We will get more into this later.

Hidden Layers

Hidden layers are the bread and butter of neural network models. Take, for example, the network pictured below with 4 input nodes and 2 hidden nodes.




Figure 1: Example of a Simple Neural Network with 4 Input Nodes and 2 Hidden Nodes

Mathematically, we can write the *activation* of the 1st node of the hidden layer, A_1 , as

$$A_1 = g(w_{0,1} + \sum_{j=1}^p w_{j,1} \cdot X_j)$$

Where $g(\cdot)$ is the nonlinear activation function of choice and $w_{j,1}$ is the weight associated with activation 1 and input node j . The value of $w_{0,1}$ is called the “bias” and can be added to offset the activation so that the minimum value matches what it is expected to be. For each activation of the hidden layer, we are taking a weighted sum of all nodes in the input layer. The activation function restricts the range of the values the activation can hold. For Sigmoid, it would be between 0 and 1 while ReLU would just be greater than or equal to 0. This can be generalized further for k activations

$$A_k = g(w_{0,k} + \sum_{j=1}^p w_{j,k} \cdot X_j)$$

Output Layer

The output layer is what we would be predicting. For a quantitative response, we would have a single node that would hold the value we predict based on the input vector X . For a categorical response with q levels, we would have q output nodes. The output can be thought of as a linear regression model fit using the hidden layer nodes as inputs. This can be formally written as

$$f(x) = \beta_0 + \sum_{k=1}^k A_k \cdot \beta_k$$

Training Neural Networks

Unlike regression models, coefficients used to predict values are not generated through Maximum Likelihood Estimation. Instead, the weights ($w_{i,j}$) are chosen through a process called *backpropagation*. Before we get into that, we should go over some other components of network training.

A popular example is a model that predicts what digit is drawn in a grid, example pictured below (3Blue1Brown 2017).




Figure 2: Example of Predicting a Number from a Hand Drawn Digit

If we input a drawing of a 3, but our network predicts it to be a 3, 5, or 9 (those output nodes would have the highest activations), we want to be able to “correct” the network. This can be done by using a *cost function*. A cost function sums up the squared differences in the activations of the nodes with what we want them to be. If we want the prediction to be a 3, we want $Y_3 = 1.00$ and $Y_1 = Y_2 = Y_4 = \dots = Y_{10} = 0.00$. Lets say Y_3 , Y_5 , and Y_9 , the output nodes associated with 3, 5, and 9, respectively, are actually equal to 0.4, 0.3, and 0.3, while the rest are close to 0. We would take the sum of the differences squared, resulting in the cost of that prediction.

Digit	Predicted	Actual	Cost
1	0.01	0	$(0 - 0.01)^2$
2	0.02	0	$(0 - 0.02)^2$
3	0.40	1	$(1 - 0.40)^2$
4	0.01	0	$(0 - 0.01)^2$
5	0.30	0	$(0 - 0.30)^2$
6	0.01	0	$(0 - 0.01)^2$
7	0.02	0	$(0 - 0.02)^2$
8	0.03	0	$(0 - 0.03)^2$
9	0.30	0	$(0 - 0.30)^2$
0	0.01	0	$(0 - 0.01)^2$

Once we know the cost of a prediction, we can calculate the change in the cost with respect to the change in any activation A or weight w using the chain rule. For example, if we were interested in how the cost C changes based on changes in the weight w , we would get

$$\frac{\delta C}{\delta w} = \frac{\delta z}{\delta w} \cdot \frac{\delta A}{\delta z} \cdot \frac{\delta C}{\delta A}$$

Where z is the weighted linear combination of inputs nodes, such that $A = g(z)$. In order to efficiently adjust the weights in the network, the algorithm known as backpropagation is used to find the gradient of the cost ∇C , which then allows us to find the negative gradient $-\nabla C$, which will tell us which weights need to be changed in order to move in the direction of decreasing cost. This process repeats until we get a minimum.

Applications of Principal Component Analysis with Neural Networks

Neural networks can grow in complexity very quickly. Given a dataset with 20 input variables, we could end up requiring many nodes in the Hidden Layer. The training process can be timely and computationally expensive.

Principal Component Analysis would allow for those 20 input variables to be trimmed down to 2 or 3 principal components. This would also theoretically cut down on the number of nodes in the Hidden Layer, thus reducing the computational cost of fitting the model while maintaining the accuracy of the model.

Single-Layer Neural Networks (Harley)

Example

Multi-Layer Neural Networks (Natasha)

Example

References

- <https://www.datacamp.com/tutorial/neural-network-models-r>
- 3Blue1Brown. 2017. “Neural Networks.” YouTube. 2017. https://www.youtube.com/watch?v=aircArvnKk&list=PLZHQObOWTQDNU6R1_67000Dx_ZCJB-3pi.
- Abdi, Hervé, Dominique Valentin, and Betty Edelman. 1999. *Neural Networks*. 124. Sage.
- Anderson, James A. 1995. *An Introduction to Neural Networks*. MIT press.
- James, Gareth, Daniela Witten, Trevor Hastie, Robert Tibshirani, et al. 2013. *An Introduction to Statistical Learning*. Vol. 112. Springer.
- Pramoditha, Rukshan. 2022. “Using PCA to Reduce Number of Parameters in a Neural Network by 30x Times.”